


# Job Queues with Gearman

Michael Moussa



# About Me

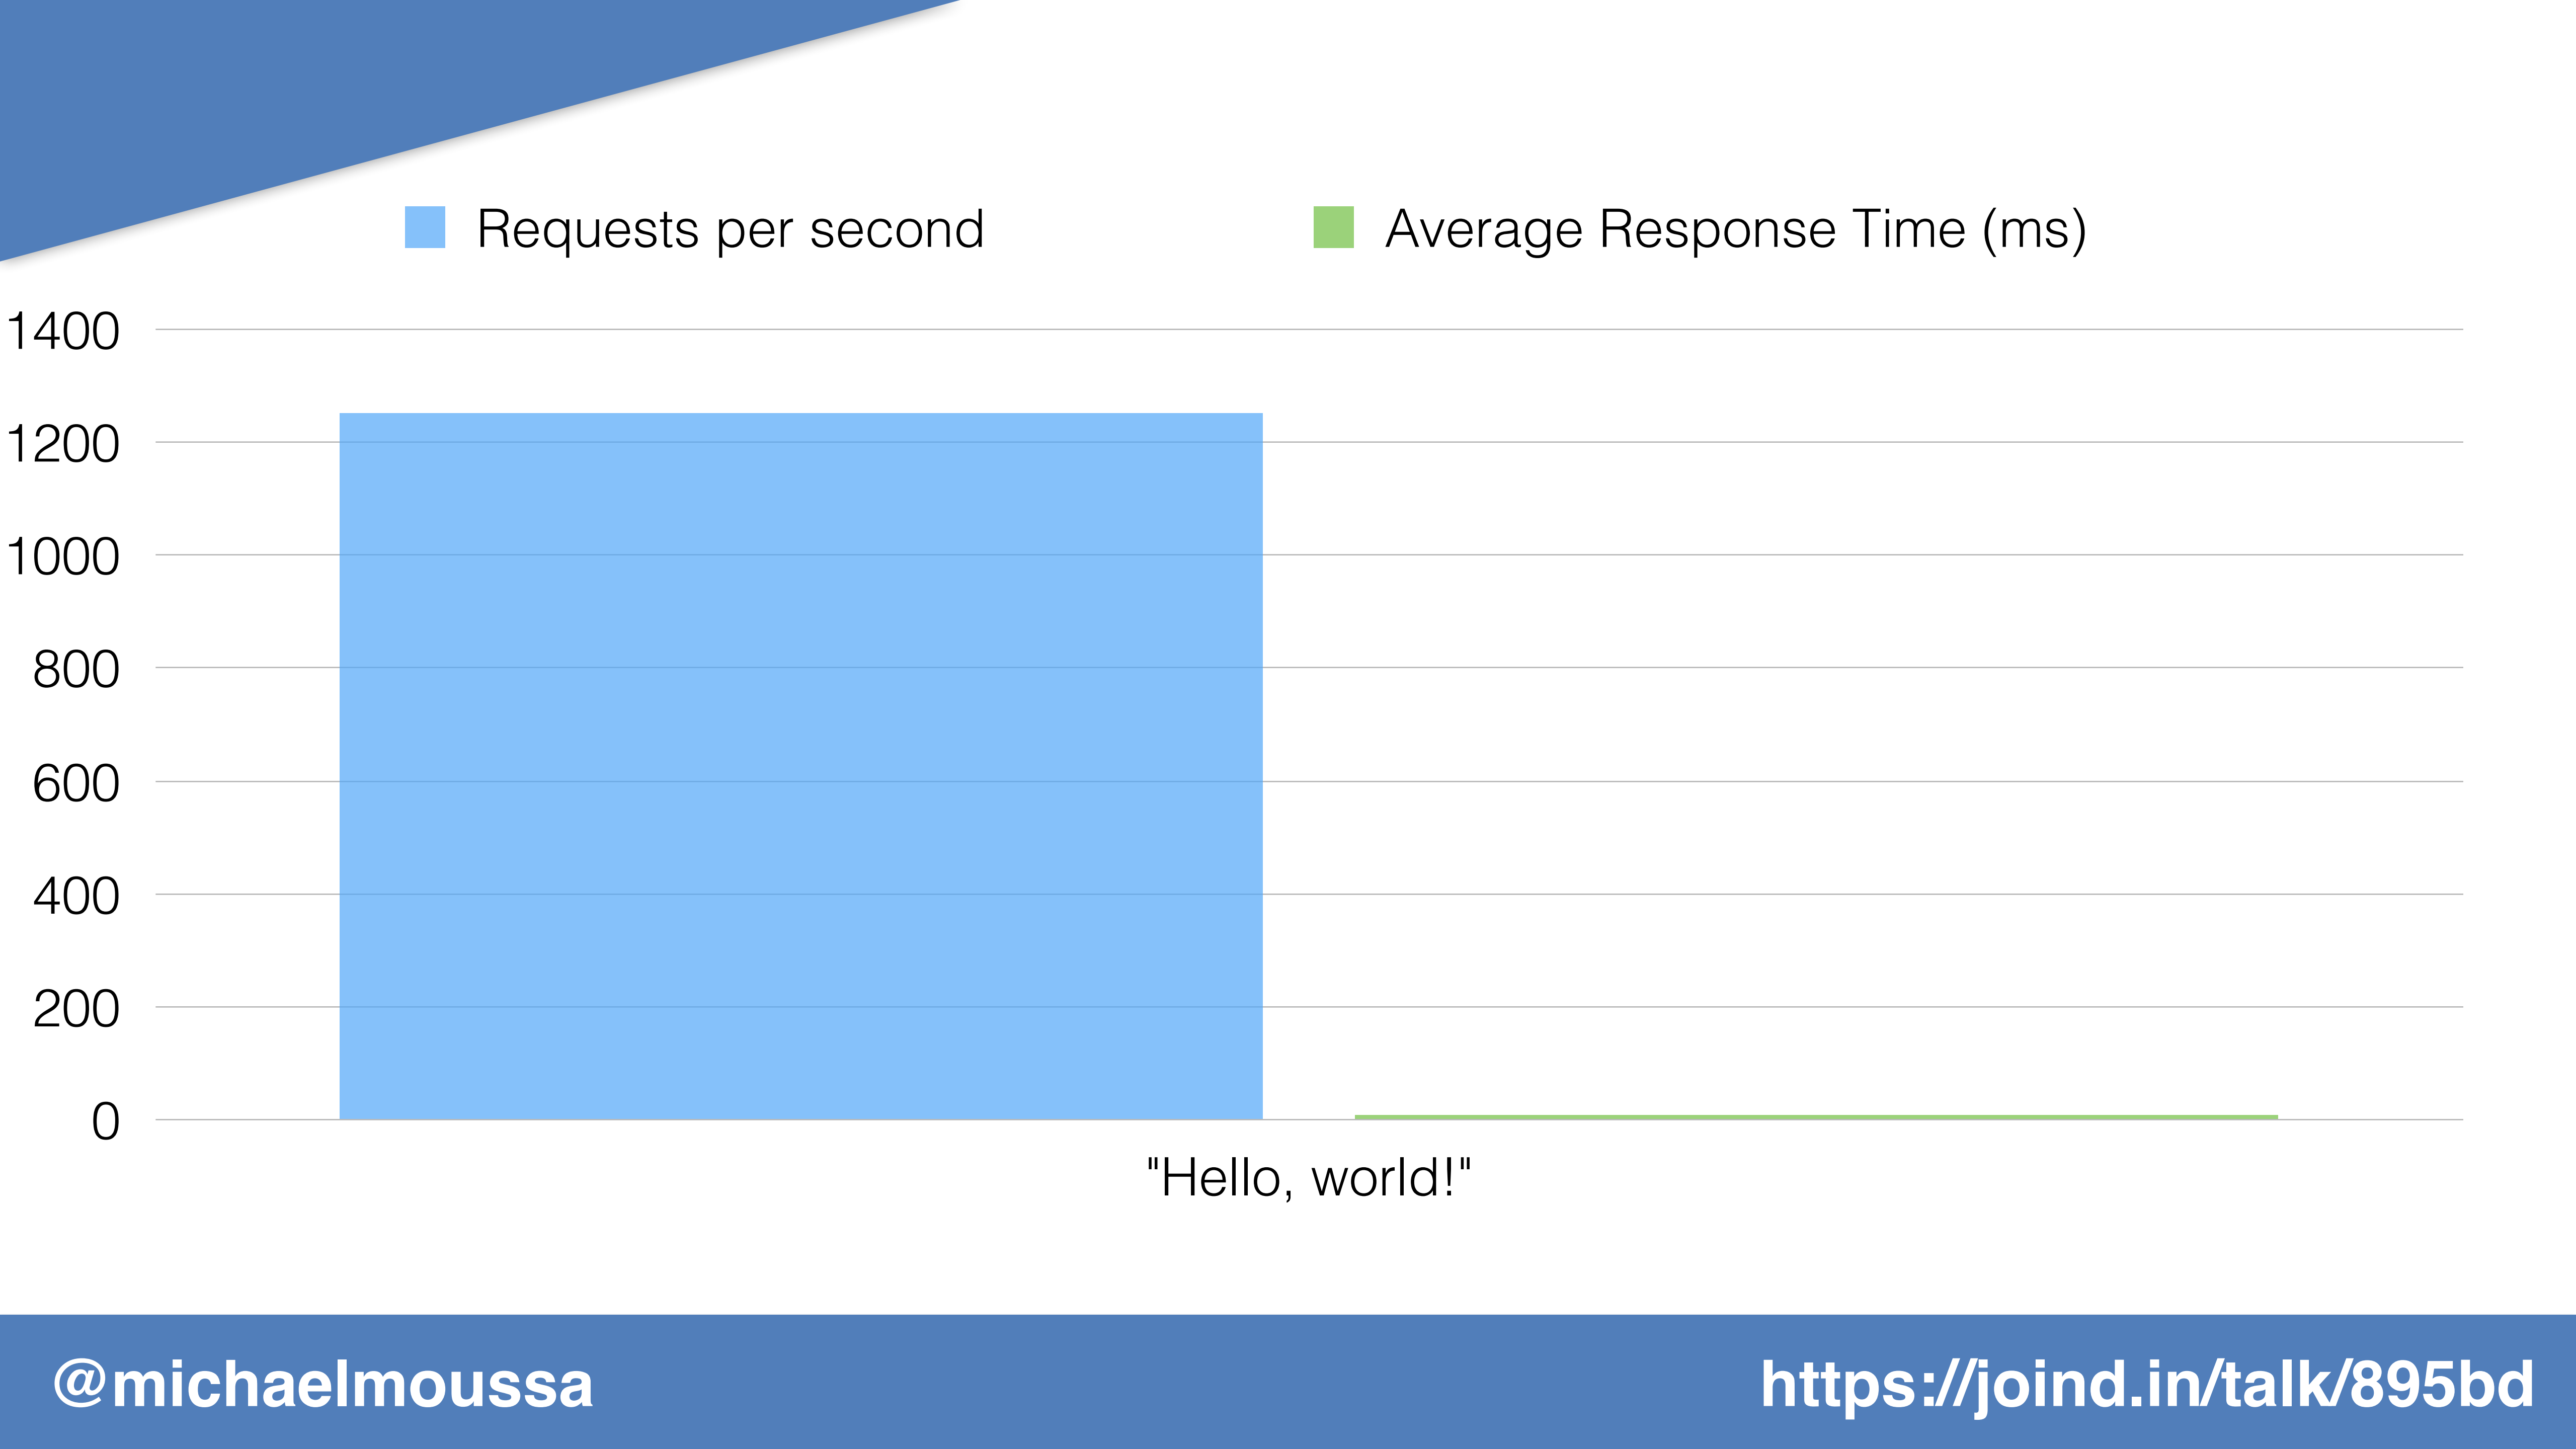
- Web application developer for 17 years
- Solutions Architect at  **rackspace**  
the #1 managed cloud company
- Zend Expressive maintainer & general open source contributor

# Why job queues?

# Benchmarks

- 1 CPU / 512MB RAM droplets running CentOS 7
- PHP 7.0, NGINX, php-fpm
- Siege (<https://github.com/JoeDog/siege>)
  - `siege -b -t 180s -c 10 http://example.com/example.php`
- Separate droplets for Siege / web / Gearman workers
- Don't worry too much about the specific numbers

```
echo 'Hello, world!';
```




```
$client = new \GuzzleHttp\Client();  
$response = $client->post(  
    'https://api.example.com/something',  
    ['json' => ['foo' => 'bar']]  
);  
  
echo (string) $response->getBody();
```



■ Requests per second

■ Average Response Time (ms)



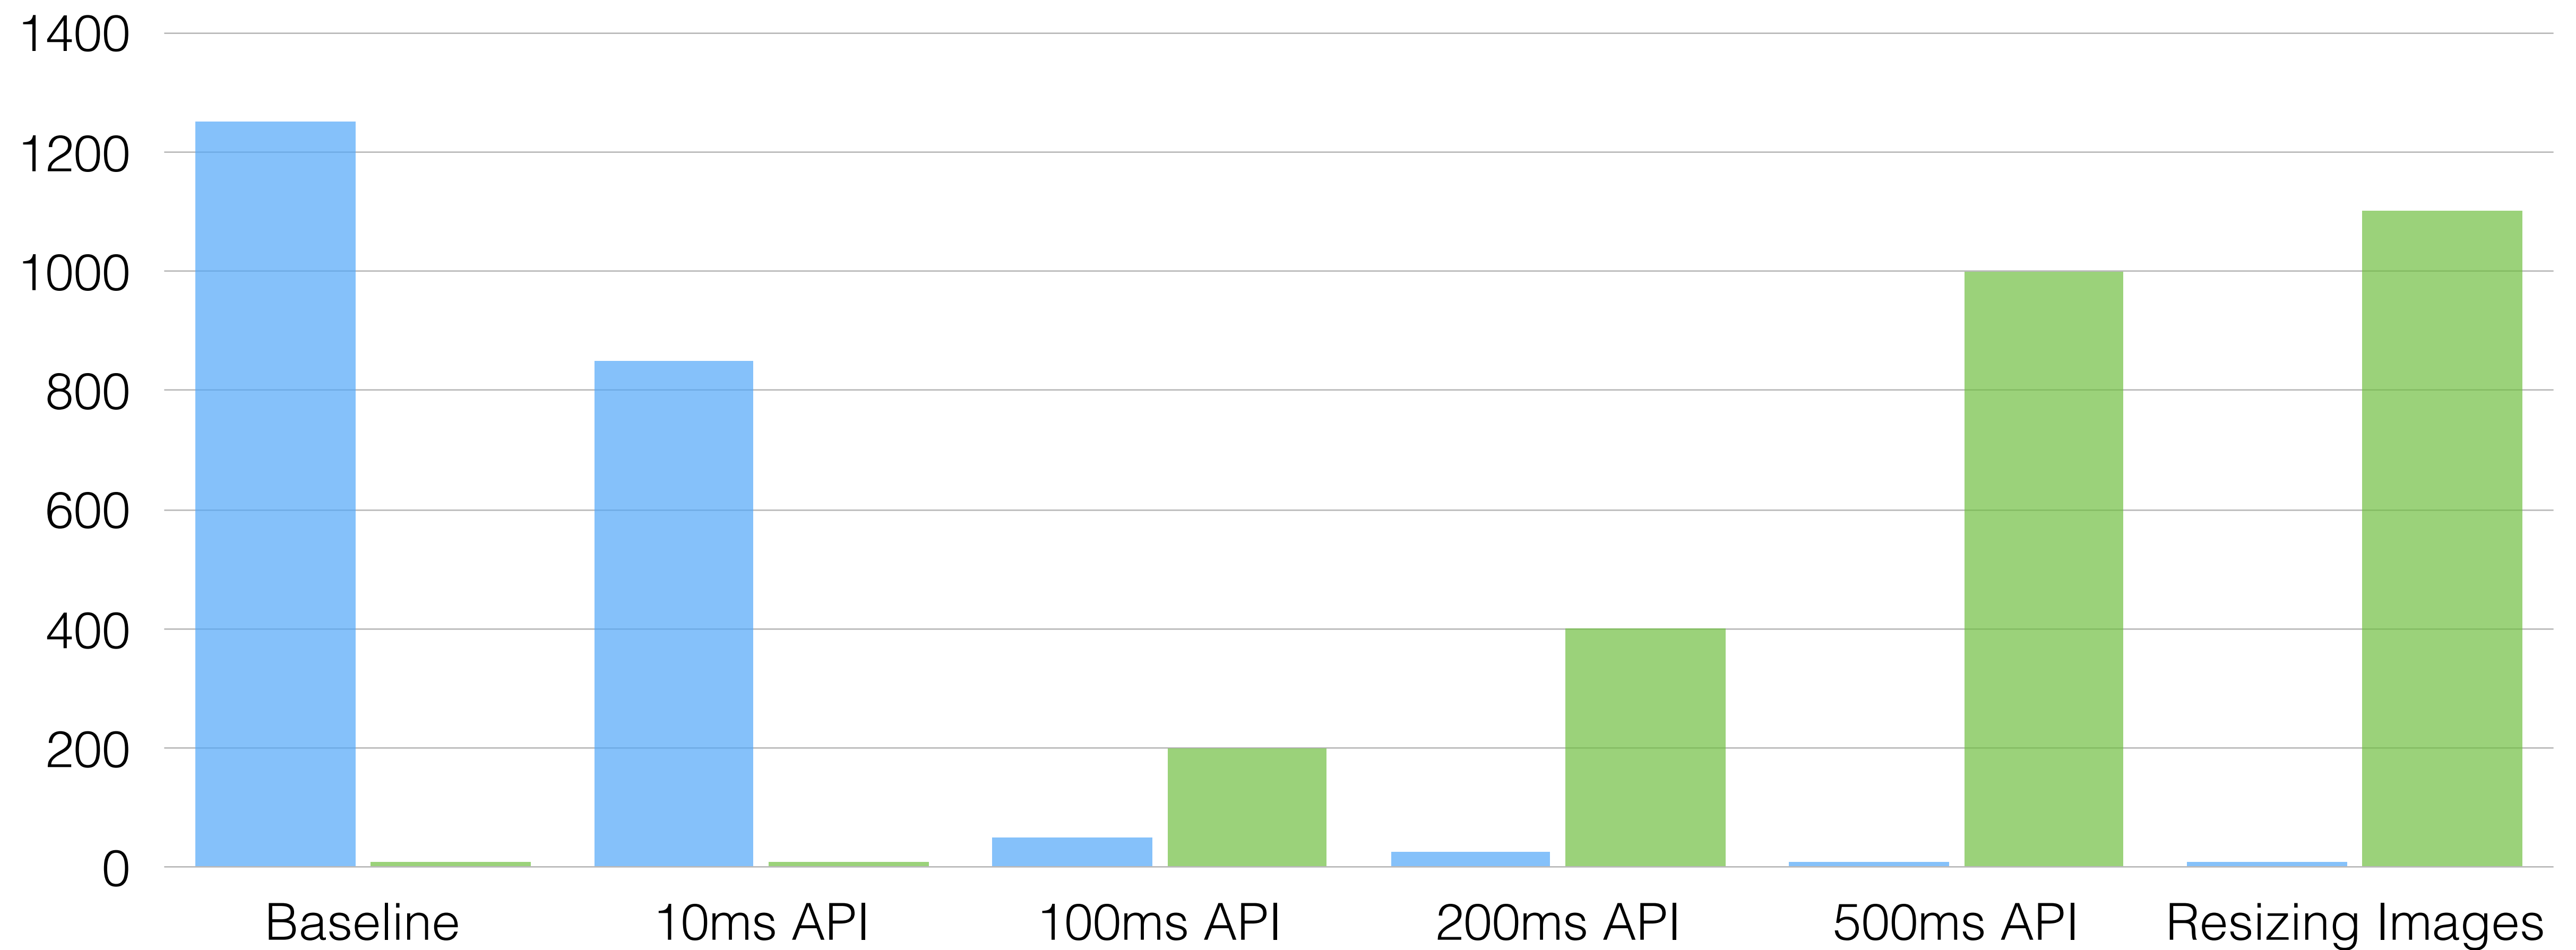


```
$imagePath = '/path/to/some-image.jpg';  
list($width, $height) = getimagesize($imagePath);  
$newWidth = $width * 0.5;  
$newHeight = $height * 0.5;  
$resizedImage = imagecreatetruecolor($newWidth, $newHeight);  
  
imagecopyresized(  
    $resizedImage, imagecreatefromjpeg($imagePath), 0, 0, 0, 0,  
    $newWidth, $newWidth, $width, $height  
);  
  
imagejpeg($resizedImage, tempnam(__DIR__, '/resized/', ''));
```



■ Requests per second

■ Average Response Time (ms)






"Gearman is a generic application framework for farming out work to other machines or processes."

<https://php.net/manual/en/intro.gearman.php>

<http://gearman.org/getting-started/>



- Workers don't read jobs from the server - the server assigns them to workers
  - No duplicate / simultaneous handling of the same job
  - Queue knows when a job is done, so you don't have to remove it
- Your worker(s) don't have to be written in PHP
- Jobs can be synchronous or asynchronous
- Jobs can be a group of tasks meant to be run in parallel

# Installation

## Gearman Server

**Ubuntu:** `apt-get install gearman-job-server`

**RHEL / CentOS:** `yum install gearmand`

## PHP Extension

**PHP 5.6:** `pecl install gearman`

### **PHP 7.0 Ubuntu:**

`add-apt-repository -y ppa:ondrej/php`

`apt-get update -y`

`apt-get install -y php-gearman`

**PHP 7.0 RHEL / CentOS:** `yum install php70-php-gearman`

• <https://rpms.remirepo.net/wizard/>

# GearmanWorker

```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);

$worker->addFunction('resize_image', function (\GearmanJob $job)
{
    $imagePath = json_decode($job->workload())->imagePath;

    // <rest of original image worker code from before goes here>
});

while ($worker->work()); // never stop working!
```



# GearmanClient

```
$client = new \GearmanClient();  
$client->addServer($host, $port);  
  
$client->doNormal(  
    'resize_image',  
    json_encode([ 'imagePath' => '/path/to/some-image.jpg' ]) )  
);
```





web01

NGINX



php fpm



```
$client = new \GearmanClient();
$client->addServer($host, $port);

echo $client->doNormal(
    'resize_image',
    json_encode(['imagePath' => '/path/to/
some-image.jpg'])
);
```



gearman01

**resize\_image**

```
{"imagePath": "/path/to/some-image.jpg"}
```

**resize\_image**

```
{"imagePath": "/path/to/some-image.jpg"}
```

**resize\_image**

```
{"imagePath": "/path/to/some-image.jpg"}
```

**resize\_image**

```
{"imagePath": "/path/to/some-image.jpg"}
```

**resize\_image**

```
{"imagePath": "/path/to/some-image.jpg"}
```

**resize\_image**

```
{"imagePath": "/path/to/some-image.jpg"}
```



worker01



```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);

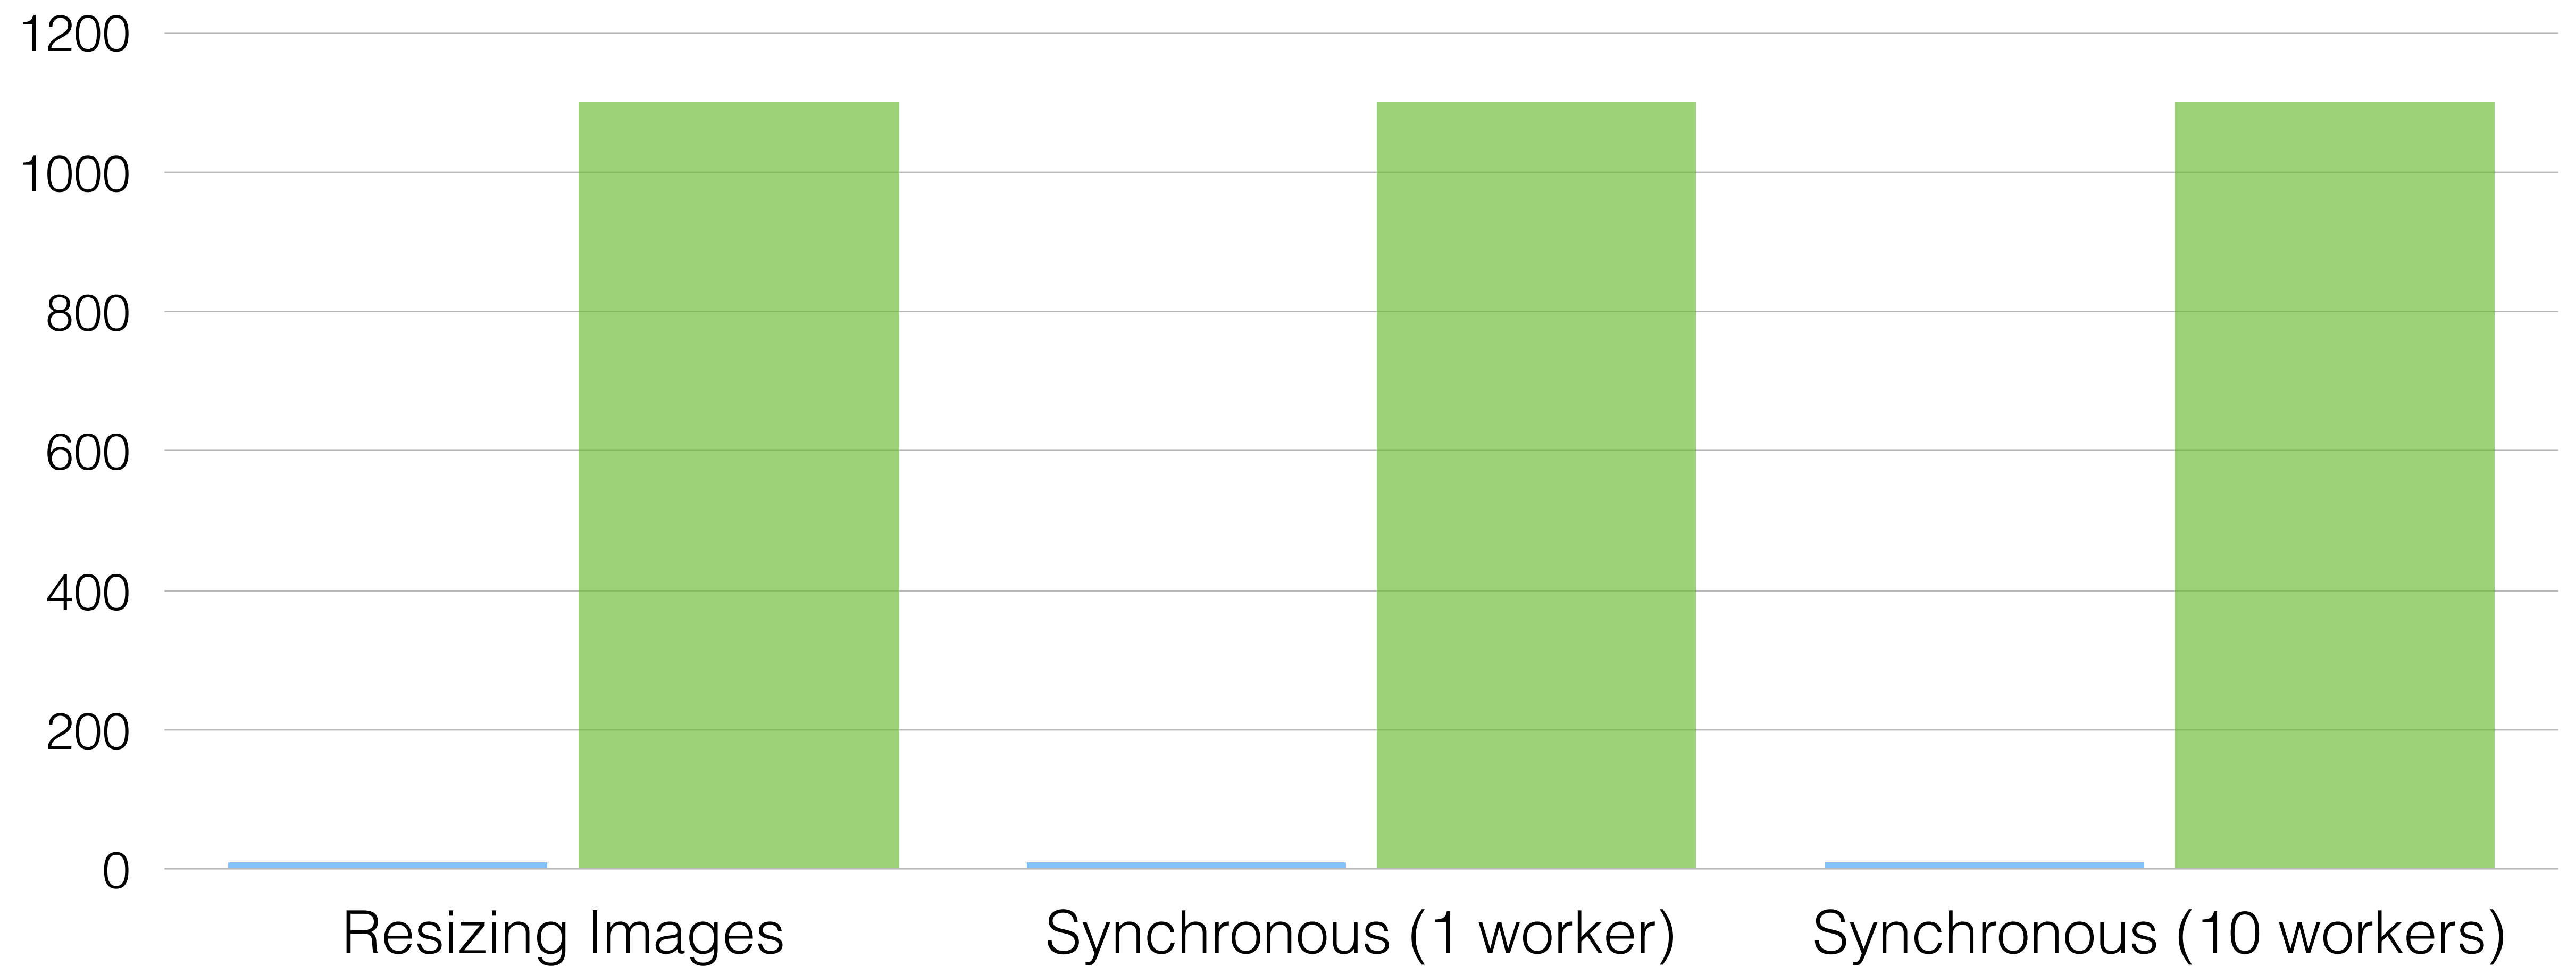
$worker->addFunction('resize_image',
function (\GearmanJob $job) {
    $imagePath = json_decode($job-
>workload())->imagePath;

    // <rest of original image worker code
from before goes here>
});

while ($worker->work()); // never stop
working!
```

■ Requests per second

■ Average Response Time (ms)



# GearmanClient

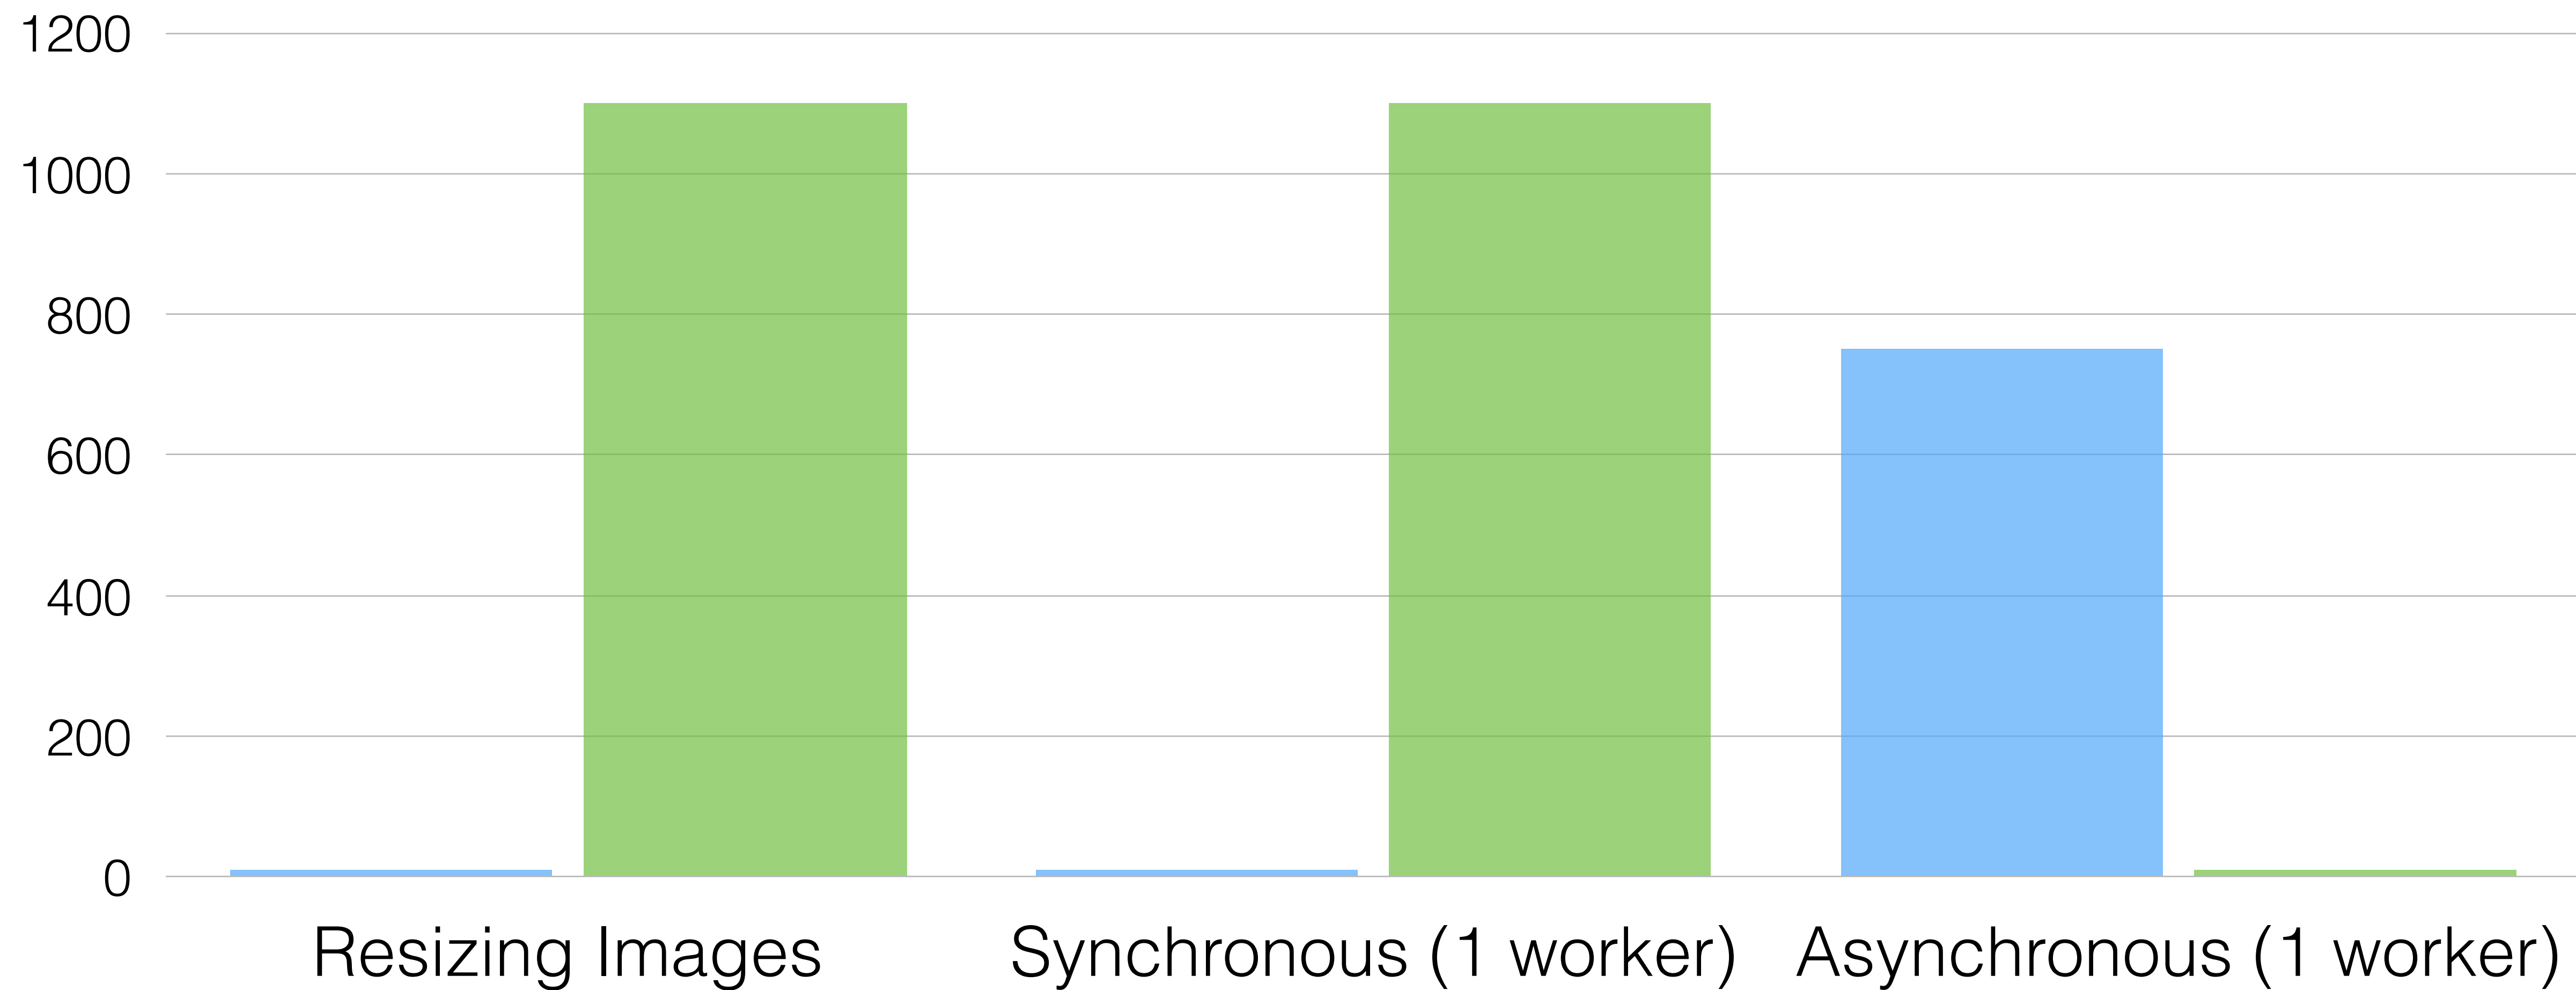
```
$client = new \GearmanClient();  
$client->addServer($host, $port);
```

```
$client->doNormal(  
$client->doBackground(  
    'resize_image',  
    json_encode([ 'imagePath' => '/path/to/some-image.jpg' ] )  
);
```



■ Requests per second

■ Average Response Time (ms)





# worker01



```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);

$worker->addFunction('resize_image',
function (\GearmanJob $job) {
    $imagePath = json_decode($job-
>workload())->imagePath;

    // <rest of original image worker code
    from before goes here>
});

while ($worker->work()); // never stop
working!
```



# worker01

php

```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);

$worker->addFunction('resize_image',
function (\GearmanJob $job) {
    $imagePath = json_decode($job-
>workload())->imagePath;

    // <rest of original image
worker code from before goes here>
});

while ($worker->work()); // never
stop working!
```

php

```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);

$worker->addFunction('resize_image',
function (\GearmanJob $job) {
    $imagePath = json_decode($job-
>workload())->imagePath;

    // <rest of original image
worker code from before goes here>
});

while ($worker->work()); // never
stop working!
```

php

```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);

$worker->addFunction('resize_image',
function (\GearmanJob $job) {
    $imagePath = json_decode($job-
>workload())->imagePath;

    // <rest of original image
worker code from before goes here>
});

while ($worker->work()); // never
stop working!
```

php

```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);

$worker->addFunction('resize_image',
function (\GearmanJob $job) {
    $imagePath = json_decode($job-
>workload())->imagePath;

    // <rest of original image
worker code from before goes here>
});

while ($worker->work()); // never
stop working!
```



# worker01

<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>
<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>
<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>
<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>	<pre>\$worker = new \GearmanWorker(); \$worker-&gt;addServer(\$host, \$port); \$worker-&gt;addFunction('re</pre>



## worker01

```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');
```



## worker02

```
$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');

$worker = new \GearmanWorker();
$worker->addServer($host, $port);
$worker->addFunction('re');
```



[illegible][illegible][illegible]

# Worker management



"Supervisor is a client/server system that allows its users to monitor and control a number of processes on UNIX-like operating systems."

<http://supervisord.org>

# Supervisor

```
# /etc/supervisord.d/resize-image-workers.ini
[program:resize-images-worker]
command=/usr/bin/php70 /path/to/image-resize-worker.php
process_name=%(program_name)s_%(process_num)02d
numprocs=5
```

---

```
$ ps -ef | grep image-resize-worker
root      2662    2661    0 14:32 ?        00:00:00 /usr/bin/php70 /path/to/image-resize-worker.php
root      2663    2661    0 14:32 ?        00:00:00 /usr/bin/php70 /path/to/image-resize-worker.php
root      2664    2661    0 14:32 ?        00:00:00 /usr/bin/php70 /path/to/image-resize-worker.php
root      2665    2661    0 14:32 ?        00:00:00 /usr/bin/php70 /path/to/image-resize-worker.php
root      2666    2661    0 14:32 ?        00:00:00 /usr/bin/php70 /path/to/image-resize-worker.php
```

# Queue management

```
$ gearadmin --status  
resize_image 100 5 5
```

- **resize\_image** - the function name
- **100** - number of jobs in queue
- **5** - number of jobs currently running
- **5** - number of capable workers

# Gearman Monitor

<https://github.com/yugene/Gearman-Monitor>



GEARMAN MONITOR

Name filter:

QueueWorkersServers

## Queue

Group by:

☒ None

☐ Server

☐ Function

Server	Function ▲	Jobs in queue	Jobs running	Workers registered
Example Gearman Server	resize_image	100	5	5

# GearmanTask

```
$client = new \GearmanClient();
$client->addServer($host, $port);

foreach ([0.25, 0.5, 0.75 as $size]) {
    $client->addTask(
        'resize_image',
        json_encode([
            'imagePath' => '/path/to/some-image.jpg',
            'size' => $size
        ])
    );
}

$client->runTasks();
```

# Task Priorities

	Foreground	Background
Low	<code>addTaskLow( ... )</code>	<code>addTaskLowBackground( ... )</code>
Normal	<code>addTask( ... )</code>	<code>addTaskBackground( ... )</code>
High	<code>addTaskHigh( ... )</code>	<code>addTaskHighBackground( ... )</code>



# Job Priorities

	Foreground	Background
Low	<code>doLow( . . . )</code>	<code>doLowBackground( . . . )</code>
Normal	<code>doNormal( . . . )</code>	<code>doBackground( . . . )</code>
High	<code>doHigh( . . . )</code>	<code>doHighBackground( . . . )</code>



# Communication

## CLIENT

```
$client->set<_____>Callback(function (\GearmanTask $task) {})
```

- Complete, Created, Data, Exception, Fail, Status, Warning, Workload

## WORKER

```
$job->send<_____>()
```

- Complete, Data, Exception, Fail, Status, Warning

```
// worker.php
```

```
$worker->addFunction(  
    'resize_image',  
    function ( \GearmanJob $job ) {  
        // do resizing here  
  
        return $job->sendComplete(  
            '/path/to/resized-image<#>'  
        );  
    }  
);
```

```
// client.php
```

```
$newImagePaths = [];
```

```
$client->setCompleteCallback(function ( \GearmanTask $task) use (&$newImagePaths) {  
    $newImagePaths[] = $task->data();  
});
```

```
$client->addTask( 'resize_image', ... );  
$client->addTask( 'resize_image', ... );  
$client->addTask( 'resize_image', ... );
```

```
$client->runTasks( );
```

```
var_dump( $newImagePaths );
```

```
// client.php
```

```
$newImagePaths = [];
```

```
$client->setCompleteCallback(function (\GearmanTask $task) use (&$newImagePaths) {  
    $newImagePaths[] = $task->data();  
});
```

```
$client->addTask('resize_image', ...);  
$client->addTask('resize_image', ...);  
$client->addTask('resize_image', ...);
```

```
$client->runTasks();
```

```
var_dump($newImagePaths);
```

```
array(3) {  
    [0]=>  
    string(23) "/path/to/resized-image1"  
    [1]=>  
    string(23) "/path/to/resized-image2"  
    [2]=>  
    string(23) "/path/to/resized-image3"  
}
```

# **Tips / Tricks / Pitfalls / Etc**

# \$unique

```
$unique = md5($functionName . '|' . $workload);
```

```
$client->doNormal($functionName, $workload, $unique);
```

- No matter how many separate clients queue the same workload, the **\$unique** identifier ensures there will only be one worked on at a time
- Foreground jobs will wait for the result from the currently active worker
- Background jobs will be discarded

# Persistence

# Persistence

```
CREATE TABLE `gearman_queue` (  
  `unique_key` varchar(64) DEFAULT NULL,  
  `function_name` varchar(255) DEFAULT NULL,  
  `priority` int(11) DEFAULT NULL,  
  `data` longblob,  
  `when_to_run` bigint(20) DEFAULT NULL,  
  UNIQUE KEY `unique_key` (  
    `unique_key`,  
    `function_name`  
  )  
) ENGINE=InnoDB DEFAULT CHARSET=utf8
```




# Persistence

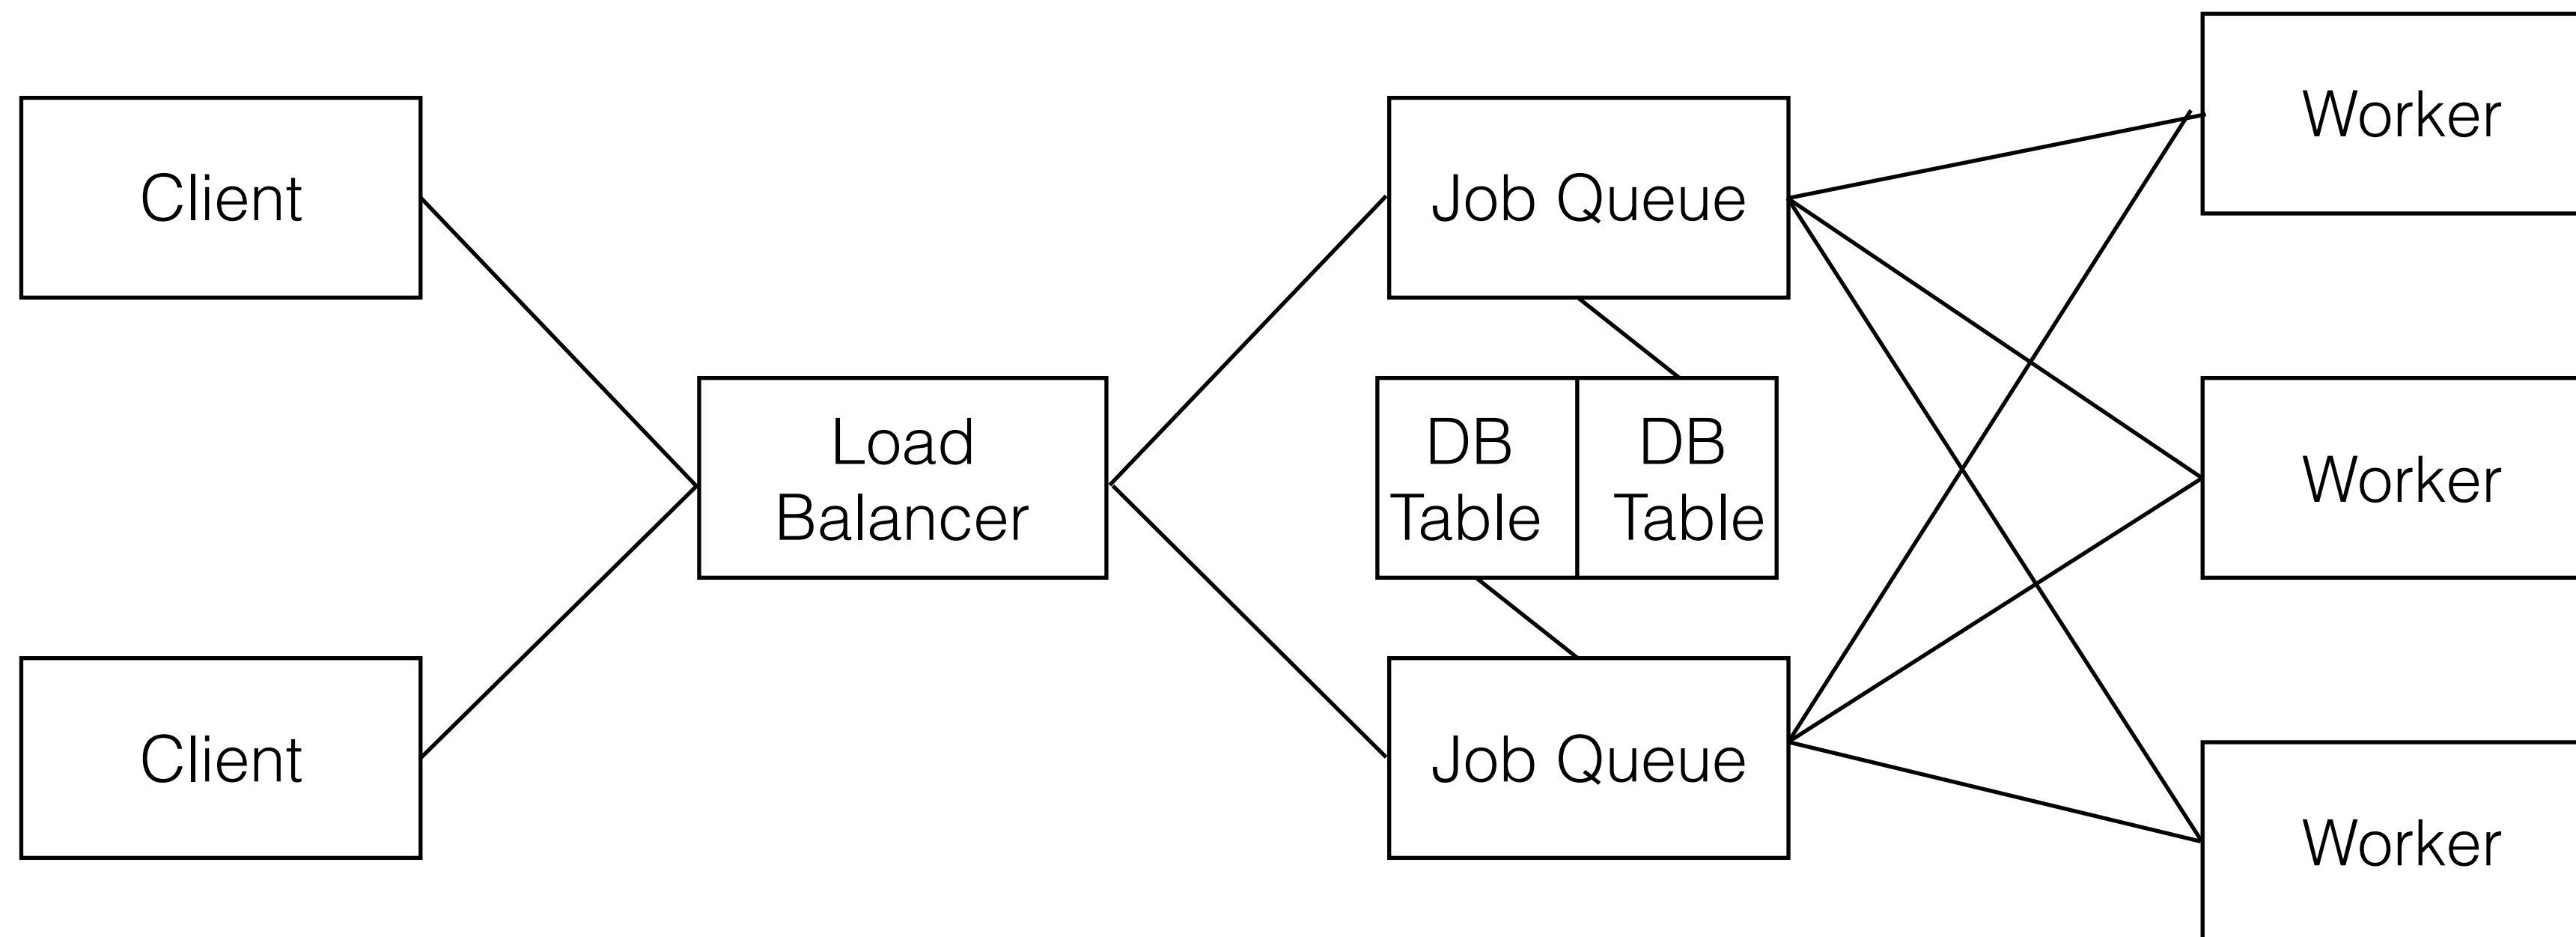
```
CREATE TABLE `gearman_queue` (  
  `unique_key` varchar(64) DEFAULT NULL,  
  `function_name` varchar(255) DEFAULT NULL,  
  `priority` int(11) DEFAULT NULL,  
  `data` longblob,  
  `when_to_run` bigint(20) DEFAULT NULL,  
  UNIQUE KEY `unique_key` (  
    `unique_key`,  
    `function_name`  
  )  
) ENGINE=InnoDB DEFAULT CHARSET=utf8
```

```
--queue-type=MySQL \  
--mysql-host=... \  
--mysql-port=3306 \  
--mysql-user=gearman \  
--mysql-password=password \  
--mysql-db=gearman \  
--mysql-table=gearman_queue
```

# High availability



# High availability



# Job size

"Gearman supports single messages up to 4GB in size!"

# Job size

~~"Gearman supports single messages up to 4GB in size!"~~

Please, no.

# Job size

~~"Gearman supports single messages up to 4GB in size!"~~

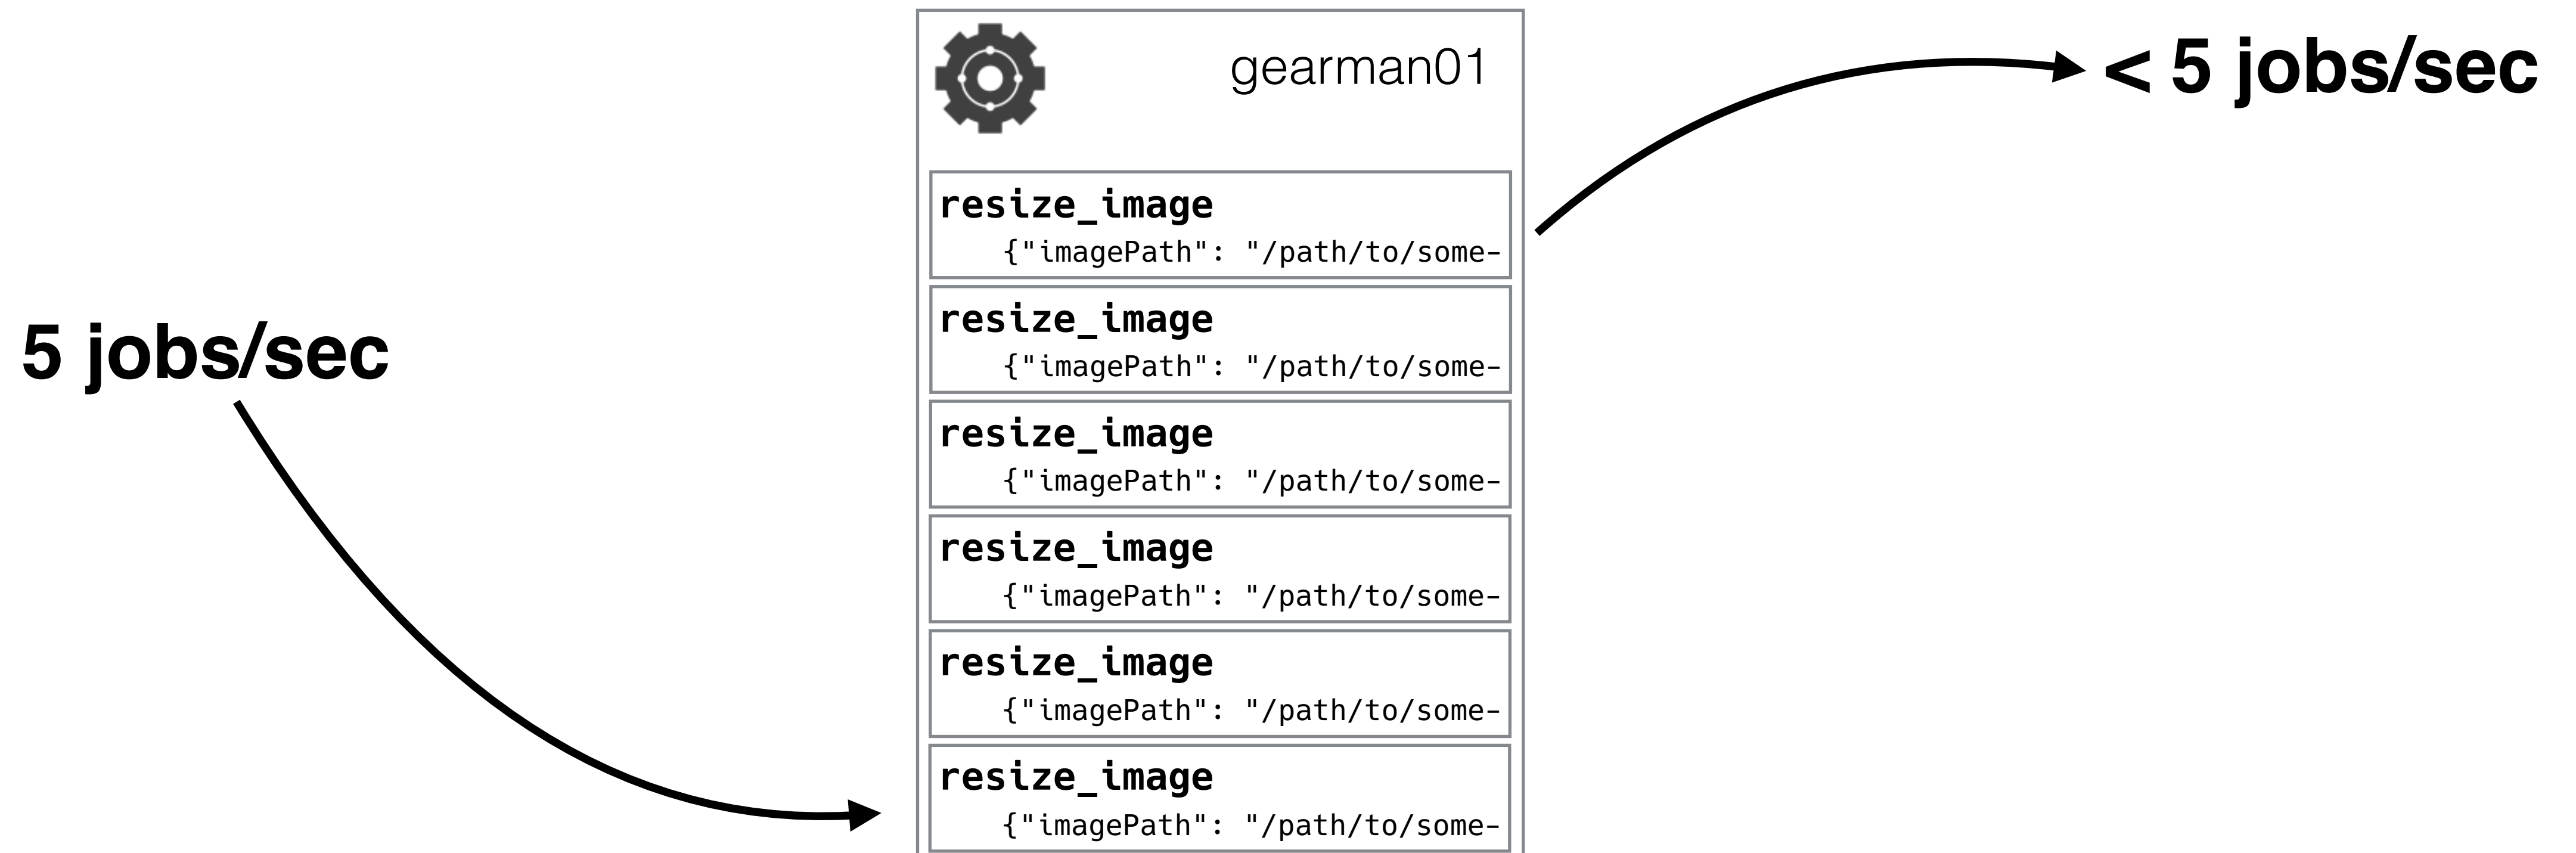
Please, no.

```
$client->doBackground(  
    'do_stuff_with_huge_file',  
    json_encode([ 's3_url' => 'https://my-bucket.s3.amazonaws.com/  
data.txt?AWSAccessKeyId=[...]&Expires=[...]&Signature=[...]' ] )  
);
```

# Capacity planning

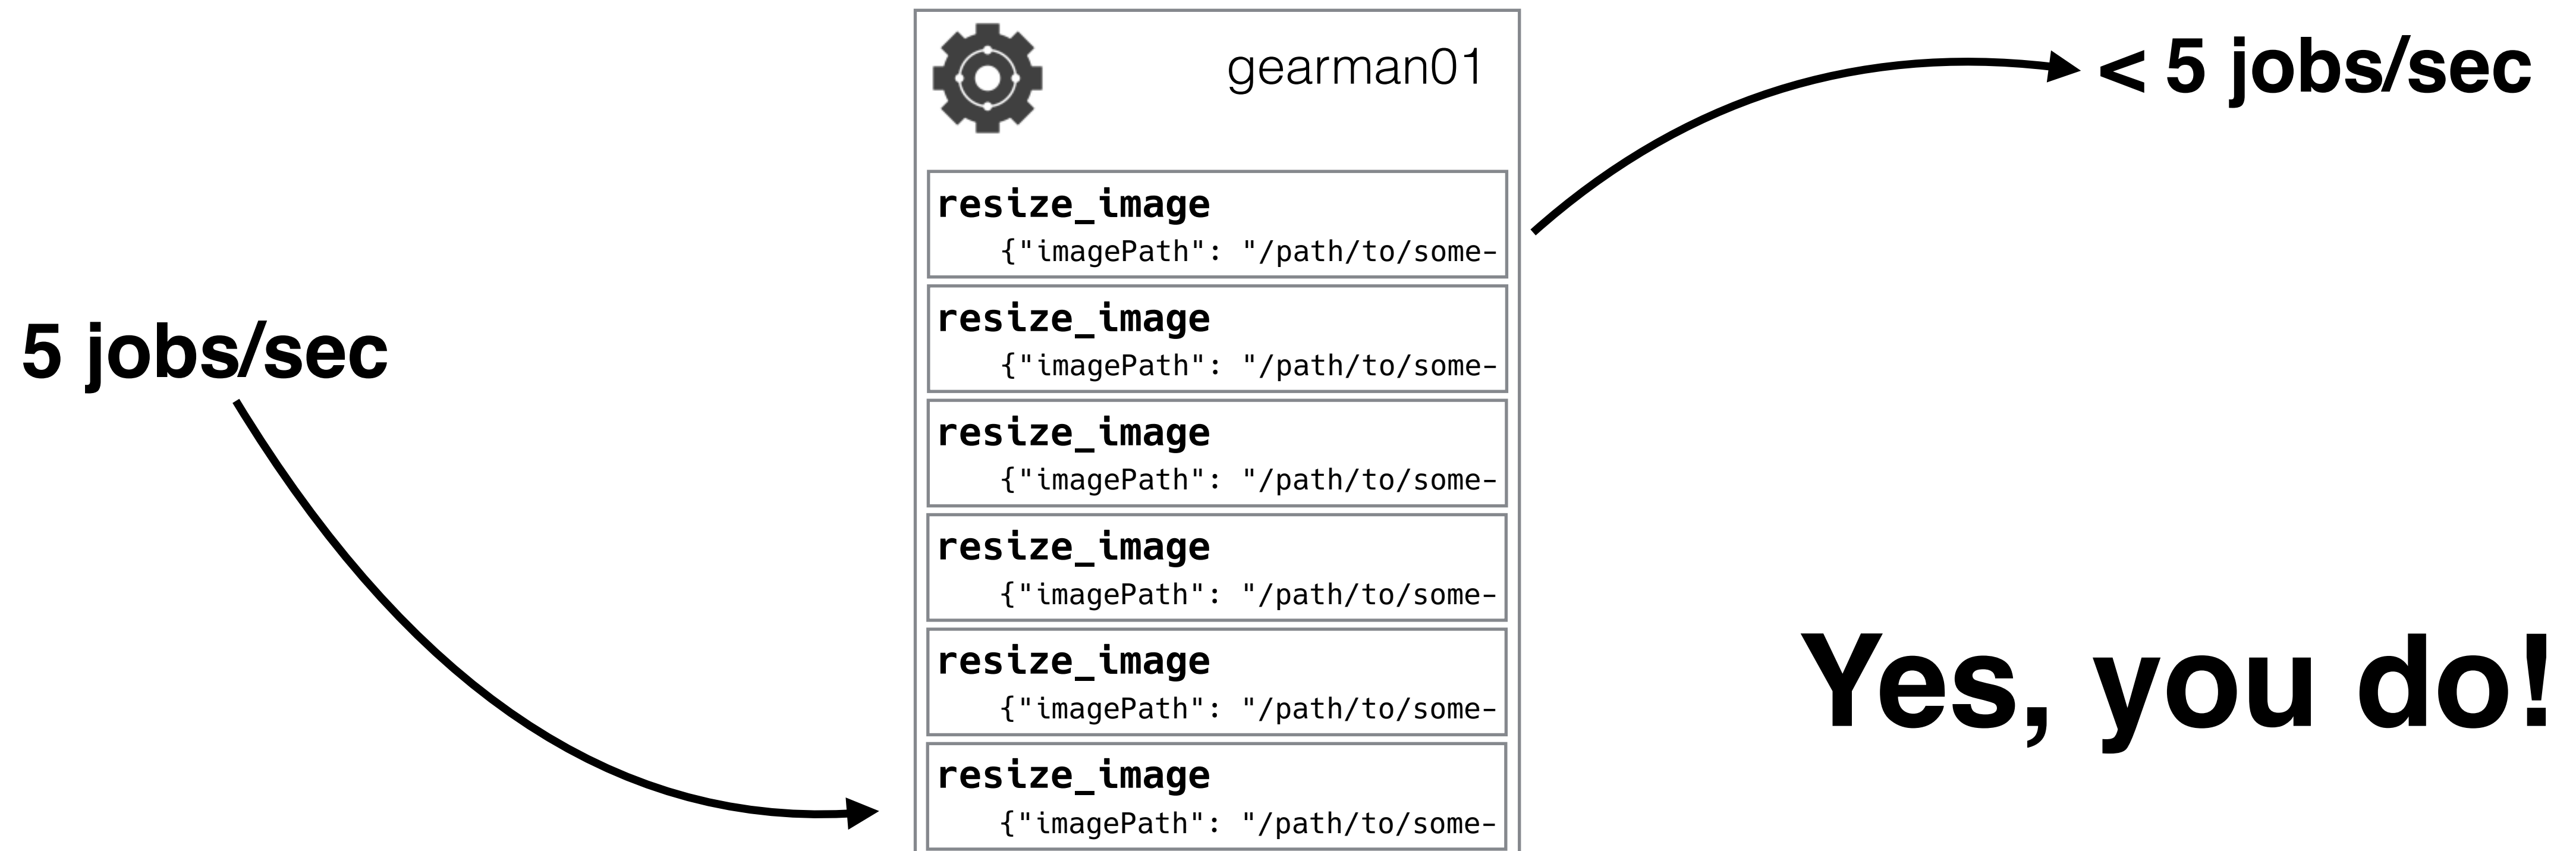
- How many worker servers do you need?
- How many workers on each one?
- How do you know when you need more?

# Do I need more?





# Do I need more?



# How many workers?

# How many workers?

"It's complicated"

# How many workers?

"It's complicated"

- Memory

# How many workers?

"It's complicated"

- Memory
- Network I/O

# How many workers?

"It's complicated"

- Memory
- Network I/O
- Disk I/O

# How many workers?

"It's complicated"

- Memory
- Network I/O
- Disk I/O
- CPU

# How many workers?

"It's complicated"

- Memory
- Network I/O
- Disk I/O
- CPU
- Load Average



# How many servers?

# How many servers?

"More than you need!"

# How many servers?

"More than you need!"

- Redundancy

# How many servers?

"More than you need!"

- Redundancy
- Peak loads

# How many servers?

"More than you need!"

- Redundancy
- Peak loads
- Unexpected new peaks

# How many servers?

"More than you need!"

- Redundancy
- Peak loads
- Unexpected new peaks
- Deployment restarts

# Who's first?

```
$client->doLowBackground( 'resize_user_avatar', '...' );  
$client->doHighBackground( 'send_welcome_email', '...' );  
$client->doBackground( 'other_stuff', '...' );
```

# Who's first?

```
$client->doLowBackground( 'resize_user_avatar', '...' );  
$client->doHighBackground( 'send_welcome_email', '...' );  
$client->doBackground( 'other_stuff', '...' );
```

**"It depends"**



# Order matters!

Competing jobs will be serviced in the order the worker functions were added

```
$worker->addFunction( 'resize_user_avatar', '...' );  
$worker->addFunction( 'send_welcome_email', '...' );  
$worker->addFunction( 'other_stuff', '...' );
```

VS

```
$worker->addFunction( 'other_stuff', '...' );  
$worker->addFunction( 'send_welcome_email', '...' );  
$worker->addFunction( 'resize_user_avatar', '...' );
```

# How about now?

```
// worker.php
```

```
$worker->addFunction( 'one', '...' );  
$worker->addFunction( 'two', '...' );  
$worker->addFunction( 'three', '...' );
```

```
// client.php
```

```
for ( $i = 0; $i < 5; $i++ ) {  
    $client->doBackground( 'one', '...' );  
    $client->doBackground( 'two', '...' );  
    $client->doBackground( 'three', '...' );  
}
```

# How about now?

```
// worker.php
```

```
$worker->addFunction( 'one', '...' );  
$worker->addFunction( 'two', '...' );  
$worker->addFunction( 'three', '...' );
```

```
// client.php
```

```
for ( $i = 0; $i < 5; $i++ ) {  
    $client->doBackground( 'one', '...' );  
    $client->doBackground( 'two', '...' );  
    $client->doBackground( 'three', '...' );  
}
```

```
one  
one  
one  
one  
one  
two  
two  
two  
two  
two  
two  
three  
three  
three  
three  
three
```

# Solution

- Forget about `$functionName`
- Think of it more like `$queueName`
- All jobs will end up in the same place and then get routed
- You only need to do this if your workers are servicing more than one function, and your queue receives a lot of traffic

```
// worker.php
```

```
$worker->addFunction('user_registration', function (\GearmanJob $job) {  
    $workload = json_decode($job->workload());  
  
    switch ($workload['function']) {  
        case 'resize_user_avatar':  
            return resizeUserAvatar($workload['data']);  
        case 'send_welcome_email':  
            return sendWelcomeEmail($workload['data']);  
        case 'other_stuff':  
            return otherStuff($workload['data']);  
    }  
});
```

# Now who's first?

```
// client.php
```

```
$client->doLowBackground(  
    'user_registration',  
    json_encode([ 'function' => 'resize_user_avatar', 'data' => [...]])  
);  
$client->doHighBackground(  
    'user_registration',  
    json_encode([ 'function' => 'send_welcome_email', 'data' => [...]])  
);  
$client->doBackground(  
    'user_registration',  
    json_encode([ 'function' => 'other_stuff', 'data' => [...]])  
);
```

# Recap



# Recap

- Offload non-critical or expensive tasks to another server using a queue

# Recap

- Offload non-critical or expensive tasks to another server using a queue
- Choose wisely between sync and async, depending on your use case

# Recap

- Offload non-critical or expensive tasks to another server using a queue
- Choose wisely between sync and async, depending on your use case
- Fill your worker servers with workers until they can't take anymore, then add more worker servers as needed

# Recap

- Offload non-critical or expensive tasks to another server using a queue
- Choose wisely between sync and async, depending on your use case
- Fill your worker servers with workers until they can't take anymore, then add more worker servers as needed
- Most of what we just covered, at a conceptual level, is applicable to all sorts of queueing systems

# Options!

- Amazon Simple Queue Service (SQS)
  - <https://aws.amazon.com/sqs/>
- RabbitMQ
  - <https://www.rabbitmq.com/>
- ZeroMQ
  - <http://zeromq.org/>
- beanstalkd
  - <https://kr.github.io/beanstalkd/>

<https://joind.in/talk/895bd>

(in case you missed it...)



# Questions?

# Thanks!