

Job Queues with Gearman

SoFloPHP Nov. 2014 - Boca
Michael Moussa | [@michaelmoussa](#)

<http://github.com/michaelmoussa/soflophp-gearman>

<http://joind.in/12749>

About Me

- PHP developer for 15 years
- Lead Developer, ZAM Network
 - lolking.net, wowhead.com, destinydb.com
- Zend Certified PHP Engineer & ZF2 Certified Architect

In a nutshell

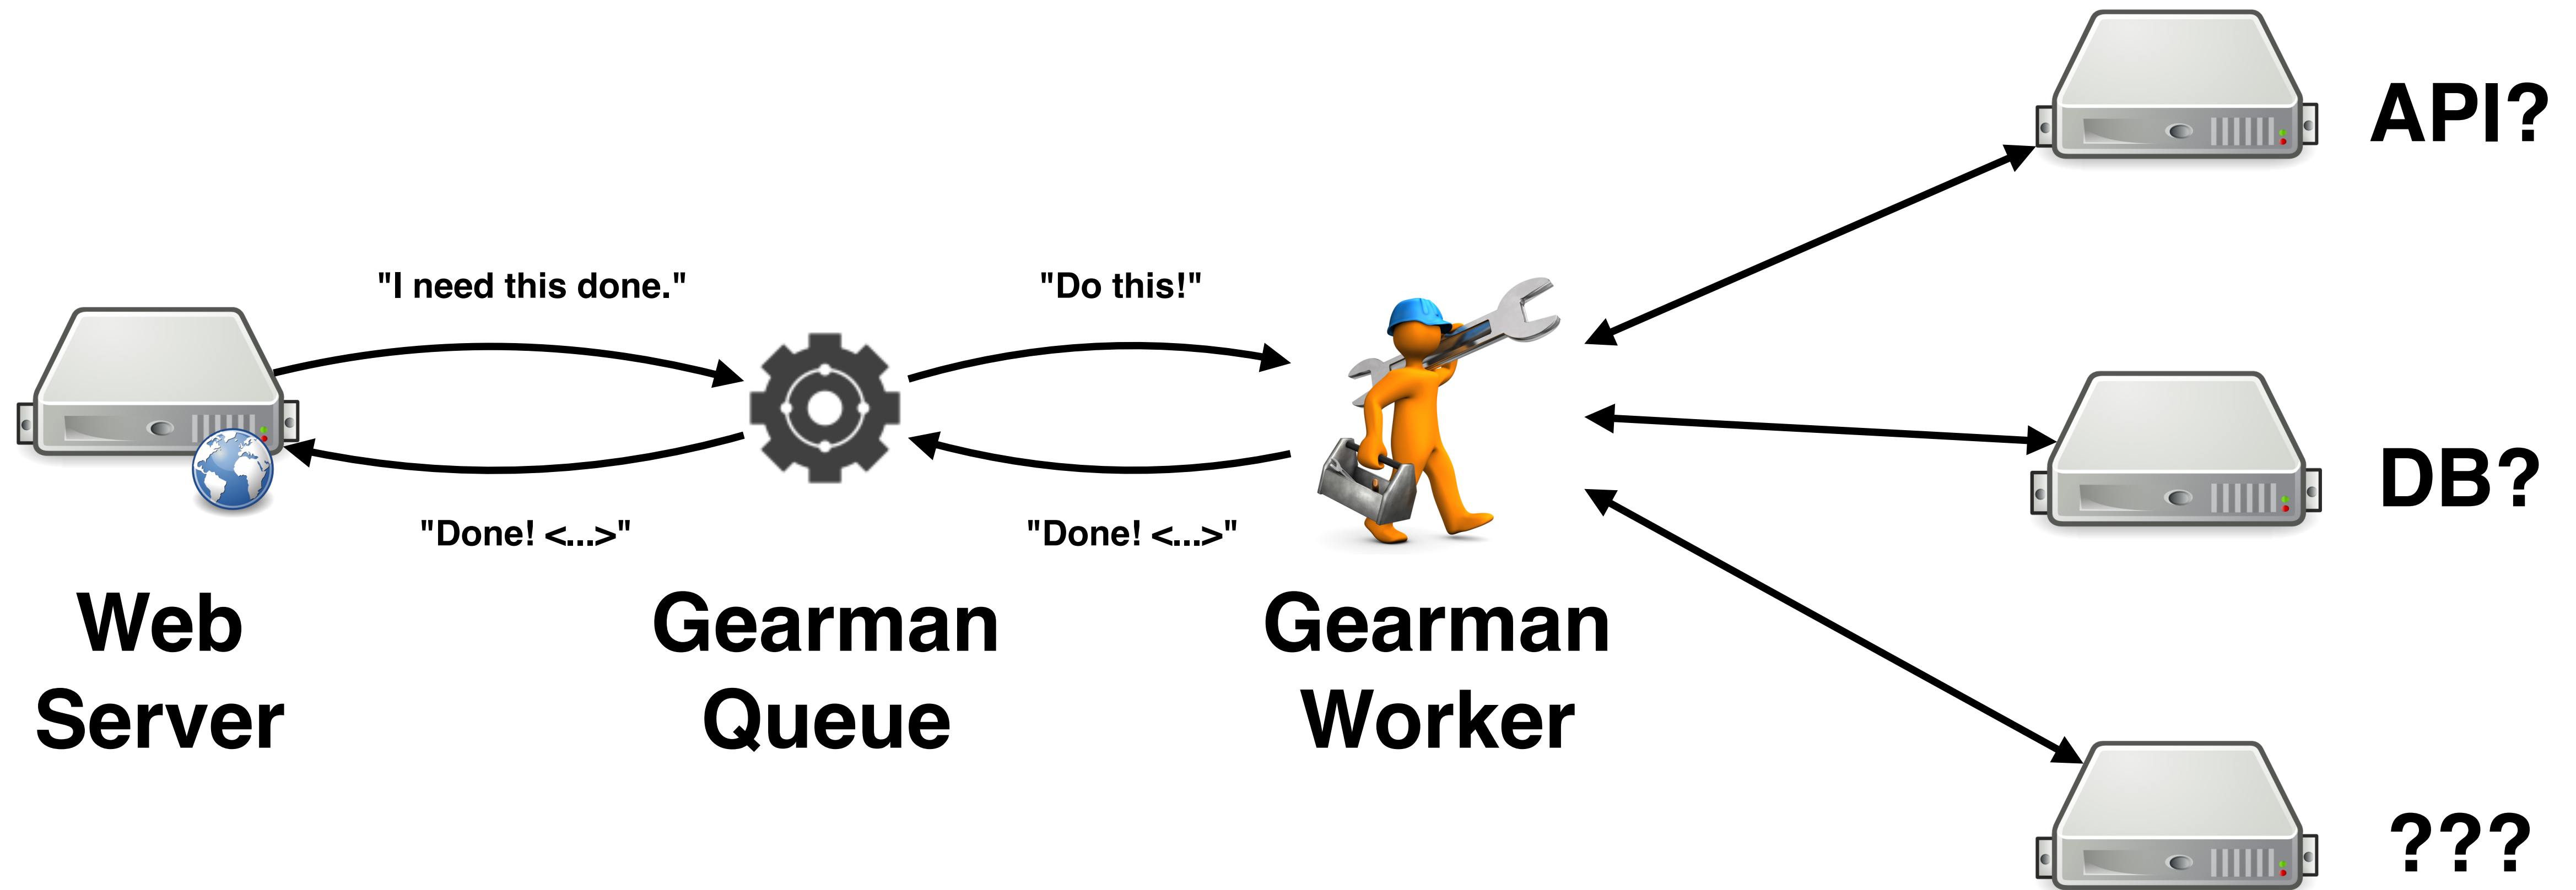
- What is Gearman?
- Why bother?
- Getting Started
- PHP API Overview
- Examples
- Unit Testing

“Gearman is a generic application framework for farming out work to multiple machines or processes.”

–<http://php.net/manual/en/intro.gearman.php>

But what does that mean?

Brief Overview



**OK... but what problem
are we trying to solve?**

DEMO TIME!

Demo 1 - "Baseline"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/1.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second:      620.39 [#/sec] (mean)
```

```
Time per request:         402.971 [ms] (mean)
```

```
...
```


Don't forget!



CPU Usage: ~13%

Demo 2 - "Doing work on the web"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/2.php"
...
Benchmarking 192.168.133.71 (be patient)
...
Requests per second:      57.22 [#/sec] (mean)
Time per request:         4368.977 [ms] (mean)
...
```



The diagram consists of three red arrows pointing from a central point to the right. The top arrow points to the 'Requests per second' value (57.22). The middle arrow points to the 'Time per request' value (4368.977). The bottom arrow points to the 'CPU Usage' value (~90%). The text '~10x worse!' is placed to the right of the middle arrow, indicating that the time per request is approximately 10 times worse than the requests per second.

~10x worse!

CPU Usage: ~90%

Your boss



or sysadmin

Gearman can help me with this?

OK - where do I start?

Getting Started

Linux (Ubuntu)

```
$ apt-get install -y apache2 php5 php5-cli php5-dev \  
    libgearman-dev gearman-job-server  
$ pecl install gearman  
$ echo "extension=gearman.so" > /etc/php5/apache2/conf.d/gearman.ini  
$ echo "extension=gearman.so" > /etc/php5/cli/conf.d/gearman.ini
```

Linux (Other distros), OSX, Windows, etc...

<http://gearman.org/getting-started/>

PHP API Summary

- GearmanClient
- GearmanJob
- GearmanTask
- GearmanWorker
- GearmanException

GearmanClient

Connects to the job server and makes requests

```
$client = new \GearmanClient();  
$client->addServer($host, $port);  
  
// Send request for some work to be done at  
// Normal priority.  
$client->doNormal('foo', '{"hello": "world"}');
```

GearmanJob

Used to represent a unit of work that needs to be done, and to communicate back to the Gearman server.

```
// Grab the "workload", i.e. parameters, and do  
// some work with them  
$result = doSomethingWith($job->workload());  
  
// Send back some indication of what happened  
$job->send<____>(...);
```

GearmanTask

Used in the client when receiving communication from server via a callback.

```
$client->setCompleteCallback(  
    function (\GearmanTask $task) {  
        doSomethingWith($task->data());  
    }  
);  
$client->addTask(...);  
$client->addTask(...);  
$client->addTask(...);  
$client->runTasks();
```



GearmanWorker

Used to register work functions and wait for jobs.

```
$worker = new \Worker();  
$worker->addServer($host, $port);  
  
// Grab "foo" jobs and call "bar" with them.  
$worker->addFunction('foo', 'bar');  
  
// Never stop working!  
while ($worker->work());
```

GearmanException

Used when something goes horribly wrong.



Don't worry... that won't happen to you.

**Now that we've got
all of that straight...**

DEMO TIME!

Demo 3 - "Foreground worker"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/3.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second: 0.47 [#/sec] (mean)
```

```
Time per request: 526542.888 [ms] (mean)
```

```
...
```

But look, web CPU is OK!

MANY times worse!

**I stopped after
~30 requests**

Web CPU Usage: ~1%

Worker CPU Usage: ~99%

Demo 4 - "Foreground workerS"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/4.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second:      23.65 [#/sec] (mean)
```

```
Time per request:         10569.333 [ms] (mean)
```

```
...
```

Web CPU still OK.

**Better than
demo #3, but
still terrible.**

Web CPU Usage: ~1%

Worker CPU Usage: ~99%

Demo 5 - "Background workerS"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/5.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second:      569.42 [#/sec] (mean)
```

```
Time per request:         439.044 [ms] (mean)
```

```
...
```

**Hey look, roughly
the same as our
baseline!**

Web CPU is fine.

Web CPU Usage: ~13%

Worker CPU Usage: ~99%

We'll talk about this later.

Quick Recap

- Don't have your web application do expensive work unless the site *requires* it and the user *expects* it to be slow.
- Run multiple workers to distribute work.
 - <http://supervisord.org/>
- Use background jobs whenever possible.
- Gearman is a tool, not magic. Be smart!

Wait, what about that high worker CPU?

- It's not *ideal*, but don't worry too much about CPU alone.
- Concern yourself more with *load average*.
 - <http://blog.scoutapp.com/articles/2009/07/31/understanding-load-averages>
- Scale out by adding more worker servers if necessary.
- It was just a demo. Your site's traffic will be much more reasonable.
- Ask your sysadmin

Tips, Tricks, and Caveats

- Use the `$uniqueId` parameter when possible.
 - e.g. `md5($functionName . $workload);`
 - Prevents queueing if an identical job is pending
- Prioritization does not work intuitively.
 - Only takes effect if `functionName` is the same.
 - Forget about "functionName". Pretend it doesn't exist.
Think of it as **QUEUE** name!
- Don't forget - the worker and client talk to each other via the server.
 - Keep your messages light. Don't give the Gearman job server more work to do than it needs.

Testing

Mocks, right?

```
$job = $this->getMock( 'GearmanJob' );  
$job->expects($this->once())  
    ->method( 'functionName' )  
    ->will($this->returnValue( 'foobar' ));  
$job->expects($this->once())  
    ->method( 'workload' )  
    ->will($this->returnValue( '{"foo": "bar"}' ));
```

No. Please don't do that.

Consider this instead

```
class TestGearmanJob extends \GearmanJob
{
    public $functionName;
    public $workload;

    public function functionName()
    {
        return $this->functionName;
    }

    public function workload()
    {
        return $this->workload;
    }
}
```

```
$job = new TestGearmanJob();
$job->functionName = 'foobar';
$job->workload = '{"foo": "bar"}';
```

Create once in
your test directory



Now you can avoid
a ton of copy & paste



Don't forget!

<https://joind.in/talk/view/12749>

So... any questions?