

# Job Queues with Gearman

SoFloPHP Nov. 2014

Michael Moussa | [@michaelmoussa](#)

<http://github.com/michaelmoussa/soflophp-gearman>

**Boca:** <https://joind.in/12749>

**Miami:** <https://joind.in/12750>

# About Me

- PHP developer for 15 years
- Lead Developer, [ZAM Network](#)
  - [lolking.net](#), [wowhead.com](#), [destinydb.com](#)
- Zend Certified PHP Engineer & ZF2 Certified Architect

# In a nutshell

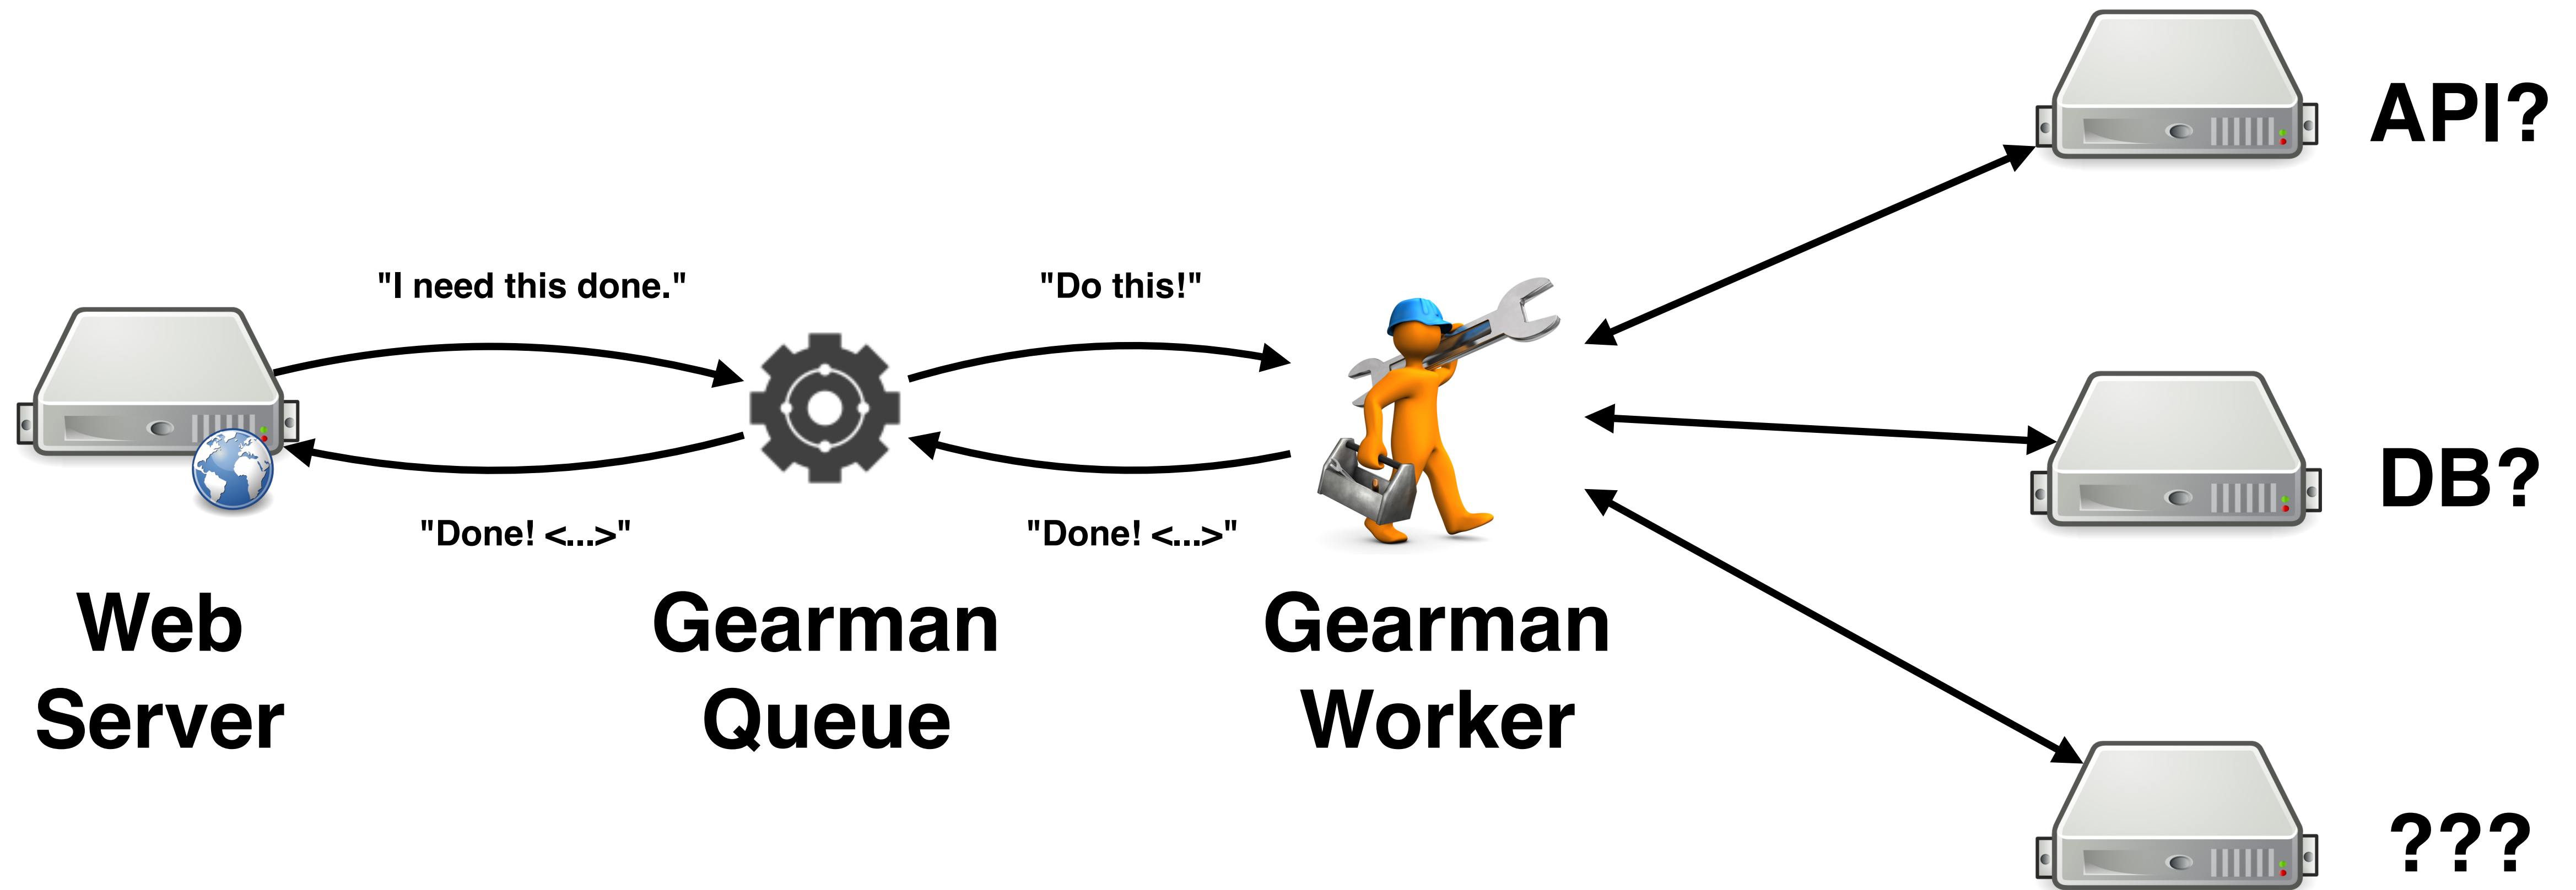
- What is Gearman?
- Why bother?
- Getting Started
- PHP API Overview
- Examples
- Unit Testing

“Gearman is a generic application framework for farming out work to multiple machines or processes.”

–<http://php.net/manual/en/intro.gearman.php>

**But what does that mean?**

# Brief Overview



**OK... but what problem  
are we trying to solve?**

DEMO TIME!

# Demo 1 - "Baseline"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/1.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second:      620.39 [#/sec] (mean)
```

```
Time per request:         402.971 [ms] (mean)
```

```
...
```

**Don't forget!**




**CPU Usage: ~13%**



## Demo 2 - "Doing work on the web"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/2.php"
...
Benchmarking 192.168.133.71 (be patient)
...
Requests per second:      57.22 [#/sec] (mean)
Time per request:         4368.977 [ms] (mean)
...
```



~10x worse!

The diagram shows three red arrows originating from the right side of the text '~10x worse!'. One arrow points to the 'Requests per second' value (57.22), another points to the 'Time per request' value (4368.977), and a third arrow points from the text '~10x worse!' down to the 'CPU Usage: ~90%' text in the next block.

**CPU Usage: ~90%**

# Your boss



# or sysadmin

**Gearman can help me with this?**

**OK - where do I start?**

# Getting Started

## Linux (Ubuntu)

```
$ apt-get install -y apache2 php5 php5-cli php5-dev \
    libgearman-dev gearman-job-server
$ pecl install gearman
$ echo "extension=gearman.so" > /etc/php5/apache2/conf.d/gearman.ini
$ echo "extension=gearman.so" > /etc/php5/cli/conf.d/gearman.ini
```

## Linux (Other distros), OSX, Windows, etc...

<http://gearman.org/getting-started/>

# PHP API Summary

- GearmanClient
- GearmanJob
- GearmanTask
- GearmanWorker
- GearmanException

# GearmanClient

Connects to the job server and makes requests

```
$client = new \GearmanClient();  
$client->addServer($host, $port);  
  
// Send request for some work to be done at  
// Normal priority.  
$client->doNormal('foo', '{"hello": "world"}');
```

# GearmanJob

Used to represent a unit of work that needs to be done, and to communicate back to the Gearman server.

```
// Grab the "workload", i.e. parameters, and do  
// some work with them  
$result = doSomethingWith($job->workload());  
  
// Send back some indication of what happened  
$job->send<____>(...);
```

# GearmanTask

Used in the client when receiving communication from server via a callback.

```
$client->setCompleteCallback(  
    function (\GearmanTask $task) {  
        doSomethingWith($task->data());  
    }  
);  
$client->addTask(...);  
$client->addTask(...);  
$client->addTask(...);  
$client->runTasks();
```




# GearmanWorker

Used to register work functions and wait for jobs.

```
$worker = new \Worker();  
$worker->addServer($host, $port);  
  
// Grab "foo" jobs and call "bar" with them.  
$worker->addFunction('foo', 'bar');  
  
// Never stop working!  
while ($worker->work());
```

# GearmanException

Used when something goes horribly wrong.



Don't worry... that won't happen to you.

**Now that we've got  
all of that straight...**

DEMO TIME!

## Demo 3 - "Foreground worker"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/3.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second: 0.47 [#/sec] (mean)
```

```
Time per request: 526542.888 [ms] (mean)
```

```
...
```

**But look, web CPU is OK!**

**MANY times worse!**

**I stopped after  
~30 requests**

**Web CPU Usage: ~1%**

**Worker CPU Usage: ~99%**

## Demo 4 - "Foreground workerS"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/4.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second:      23.65 [#/sec] (mean)
```

```
Time per request:         10569.333 [ms] (mean)
```

```
...
```

**Web CPU still OK.**

**Better than  
demo #3, but  
still terrible.**

**Web CPU Usage: ~1%**

**Worker CPU Usage: ~99%**

## Demo 5 - "Background workerS"

```
$ ab -n 10000 -c 250 "http://192.168.133.71/demo/5.php"
```

```
...
```

```
Benchmarking 192.168.133.71 (be patient)
```

```
...
```

```
Requests per second:      569.42 [#/sec] (mean)
```

```
Time per request:         439.044 [ms] (mean)
```

```
...
```

**Web CPU is fine.**

**Hey look, roughly  
the same as our  
baseline!**

**Web CPU Usage: ~13%**

**Worker CPU Usage: ~99%**

**We'll talk about this later.**



# Quick Recap

- Don't have your web application do expensive work unless the site *requires* it and the user *expects* it to be slow.
- Run multiple workers to distribute work.
  - <http://supervisord.org/>
- Use background jobs whenever possible.
- Gearman is a tool, not magic. Be smart!

# Wait, what about that high worker CPU?

- It's not *ideal*, but don't worry too much about CPU alone.
- Concern yourself more with *load average*.
  - <http://blog.scoutapp.com/articles/2009/07/31/understanding-load-averages>
- Scale out by adding more worker servers if necessary.
- It was just a demo. Your site's traffic will be much more reasonable.
- Ask your sysadmin



# Tips, Tricks, and Caveats

- Use the `$uniqueId` parameter when possible.
  - e.g. `md5($functionName . $workload);`
  - Prevents queueing if an identical job is pending
- Prioritization does not work intuitively.
  - Only takes effect if `functionName` is the same.
  - Forget about "functionName". Pretend it doesn't exist.  
Think of it as **QUEUE** name!
- Don't forget - the worker and client talk to each other via the server.
  - Keep your messages light. Don't give the Gearman job server more work to do than it needs.

# Testing

## Mocks, right?

```
$job = $this->getMock( 'GearmanJob' );  
$job->expects($this->once())  
    ->method( 'functionName' )  
    ->will($this->returnValue( 'foobar' ));  
$job->expects($this->once())  
    ->method( 'workload' )  
    ->will($this->returnValue( '{ "foo": "bar" }' ));
```

**No. Please don't do that.**

# Consider this instead

```
class TestGearmanJob extends \GearmanJob
{
    public $functionName;
    public $workload;

    public function functionName()
    {
        return $this->functionName;
    }

    public function workload()
    {
        return $this->workload;
    }
}
```

```
$job = new TestGearmanJob();
$job->functionName = 'foobar';
$job->workload = '{"foo": "bar"}';
```

Create once in  
your test directory



Now you can avoid  
a ton of copy & paste



# Don't forget!

**Boca:** <https://joind.in/12749>

**Miami:** <https://joind.in/12750>

# So... any questions?